



ELSEVIER

Computational Statistics & Data Analysis 36 (2001) 279–298

COMPUTATIONAL
STATISTICS
& DATA ANALYSIS

www.elsevier.com/locate/csda

Bandwidth selection for kernel conditional density estimation

David M. Bashtannyk, Rob J. Hyndman *

Department of Econometrics and Business Statistics, Monash University, Clayton, Melbourne, Vic. 3800, Australia

Received 1 August 1999; accepted 1 September 2000

Abstract

We consider bandwidth selection for the kernel estimator of conditional density with one explanatory variable. Several bandwidth selection methods are derived ranging from fast rules-of-thumb which assume the underlying densities are known to relatively slow procedures which use the bootstrap. The methods are compared and a practical bandwidth selection strategy which combines the methods is proposed. The methods are compared using two simulation studies and a real data set. © 2001 Elsevier Science B.V. All rights reserved.

Keywords: Bandwidth selection; Conditioning; Density estimation; Kernel smoothing

1. Introduction

To motivate the problem, consider the data given in Azzalini and Bowman (1990) on the waiting time between the starts of successive eruptions and the duration of the subsequent eruption for the Old Faithful geyser in Yellowstone National Park, Wyoming. The data were collected continuously from August 1st until August 15th, 1985. There are a total of 299 observations. The times are measured in minutes. Some duration measurements, taken at night, were originally recorded as S (short), M (medium), and L (long). These values have been coded as 2, 3 and 4 min, respectively. This data set is also distributed with S-Plus (2000).

Fig. 1 shows a scatterplot of the data. Clearly, when there has been a relatively short waiting time between eruptions, the duration of the next eruption is relatively

* Corresponding author.

E-mail address: rob.hyndman@buseco.monash.edu.au (R.J. Hyndman).




Fig. 1. Old Faithful Geyser data: duration of eruption plotted against waiting time to the eruption.




Fig. 2. Estimated conditional density of eruption duration conditional on waiting time to the eruption. Bandwidths chosen using the combination method.

long (more than 3.5 min). However, when the waiting time between eruptions is longer than 70 min, there is a mixture of short durations (less than 2.5 min) and long durations (more than 3.5 min) with very few between. Hence, the conditional density of duration given a waiting time of 80 min is bimodal, whereas the conditional density of duration given a waiting time of 50 min is unimodal.

We are interested in estimating the conditional density using kernel methods. Some kernel conditional density estimates for the data shown in Fig. 1 are shown in Fig. 2. Note the shift in mean as well as the change in modality as waiting time changes.

Standard nonparametric regression does not allow the analysis of changes in modality, and standard density estimation does not allow conditioning on an explanatory variable. Conditional density estimation is, in some ways, a generalization of both nonparametric regression and standard univariate density estimation.

Kernel conditional density estimation was first considered by Rosenblatt (1969) who studied the problem of estimating the density of Y conditional on $X = x$ where X is univariate and random. If $g(x, y)$ denotes the joint density of (X, Y) and $h(x)$ denotes the marginal density of X , then the conditional density of $Y | (X = x)$ is given by $f(y | x) = g(x, y)/h(x)$.

Hyndman, Bashtannyk and Grunwald (1996) considered the following modified form of Rosenblatt's estimator:

$$\hat{f}(y | x) = \frac{(1/nab) \sum_{j=1}^n K((\|x - X_j\|_x)/a) K((\|y - Y_j\|_y)/b)}{(1/na) \sum_{j=1}^n K((\|x - X_j\|_x)/a)}, \quad (1.1)$$

where $\{(X_1, Y_1), \dots, (X_n, Y_n)\}$ is a sample of independent observations from the distribution of (X, Y) and $\|\cdot\|_x$ and $\|\cdot\|_y$ are distance metrics on the spaces of X and Y , respectively.

The kernel function, $K(u)$, is assumed to be a real, integrable, nonnegative, even function on \mathbb{R} concentrated at the origin such that

$$\int_{\mathbb{R}} K(u) du = 1, \quad \int_{\mathbb{R}} u K(u) du = 0 \quad \text{and} \quad \sigma_K^2 = \int_{\mathbb{R}} u^2 K(u) du < \infty. \quad (1.2)$$

Popular choices for $K(u)$ are defined in terms of univariate and unimodal probability density functions.

The problem of conditional density estimation appears to have lain free of scrutiny until it was revisited recently and some improved estimators were proposed.

Hyndman et al. (1996) give the bias, variance, MSE and convergence properties of $\hat{f}(y | x)$ and proposed an alternative kernel estimator with smaller MSE than the standard estimator in some commonly occurring situations. Fan, Yao and Tong (1996) proposed an alternative conditional density estimator by generalizing Rosenblatt's estimator using local polynomial techniques. Hyndman and Yao (1998) introduced two further local parametric estimators which improve on the estimators given by Fan et al. (1996). Stone (1994) followed a different path and considered using tensor products of polynomial splines to obtain conditional log density estimates.

In this paper we consider the problem of bandwidth selection for estimator (1.1). We also comment on how to extend the ideas presented here to the improved estimators introduced later.

We shall rewrite (1.1) as

$$\hat{f}(y | x) = \frac{1}{b} \sum_{j=1}^n w_j(x) K\left(\frac{\|y - Y_j\|_y}{b}\right) \quad \text{where} \quad w_j(x) = \frac{K((\|x - X_j\|_x)/a)}{\sum_{i=1}^n K((\|x - X_i\|_x)/a)}. \quad (1.3)$$

The parameters a and b control the degree of smoothing applied to the density estimate; a controls the smoothness between conditional densities in the x direction and b controls the smoothness of each conditional density in the y direction. The

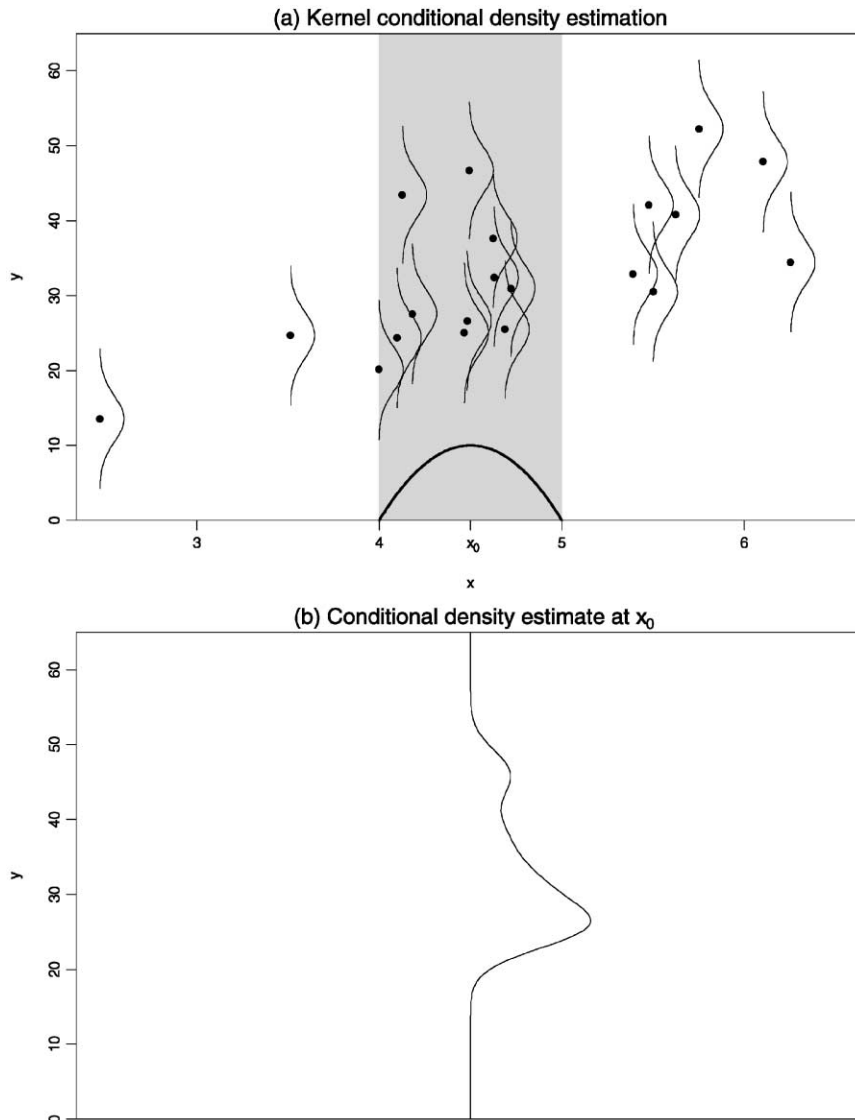


Fig. 3. Construction of the kernel conditional density estimate $f(y|x_0)$ at the conditioning value $X=x_0$. The shaded region shows the observations which receive nonzero weight. The weight function is shown as the heavy line in the top plot.

selection of a and b has a critical role in determining the performance of the kernel conditional density estimate.

Fig. 3 shows graphically how the kernel conditional density estimate is constructed. For simplicity we have used 20 observations, although a much higher number of observations is required for meaningful conditional density estimation.

Fig. 3(a) shows kernel functions with bandwidth b , centered at the observations. The conditioning $X=x$ is carried out by another kernel function in the X -space. This

second kernel function has bandwidth a and is centered at the conditioning value $x = x_0$. (This kernel is normalized so that the total weights sum to one.) The kernel function chosen for this illustration has bounded support and observations outside the window width a carry zero weight. The shaded region shows those observations which have nonzero weight.

In Fig. 3(b) the conditional density estimate at $X = x_0$ is shown. This was obtained by summing the n kernel functions in Y -space, weighted by $\{w_j(x)\}$ in X -space.

Our approach in bandwidth selection will be to minimize the weighted integrated mean square error function (IMSE), defined as (see, e.g., Wand and Jones, 1995)

$$\text{IMSE}(a, b; \hat{f}, f) = \int \int E\{\hat{f}(y|x) - f(y|x)\}^2 h(x) dx dy. \quad (1.4)$$

Weighting the IMSE by the marginal density $h(x)$ places more emphasis on the regions that have more data and it also eases the computational difficulty.

We also define the integrated square error function (ISE) as (Wand and Jones, 1995)

$$\text{ISE}(a, b; \hat{f}, f) = \int \int \{\hat{f}(y|x) - f(y|x)\}^2 h(x) dx dy. \quad (1.5)$$

Note that this is the expected value of $\int \{\hat{f}(y|X) - f(y|x)\}^2 dy$ with respect to X . For numerical examples, we will estimate the ISE using

$$I(a, b; \mathbf{X}, \mathbf{Y}, \mathbf{y}', f) = \frac{\Delta}{n} \sum_{j=1}^N \sum_{i=1}^n [\hat{f}(y'_j | X_i) - f(y'_j | X_i)]^2, \quad (1.6)$$

where $\mathbf{X} = \{X_1, \dots, X_n\}$, $\mathbf{Y} = \{Y_1, \dots, Y_n\}$ and $\{(X_i, Y_i)\}$ is an iid sample with density $g(\cdot, \cdot)$, $\mathbf{y}' = \{y'_1, \dots, y'_N\}$ is a vector of equally spaced values over the sample space of Y with $y_{i+1} - y_i = \Delta$, and \hat{f} is calculated from $\{(X_i, Y_i)\}$ using (1.3). We average (1.6) across samples to estimate the IMSE using

$$\hat{M}(a, b; m, \mathbf{y}', f) = \frac{1}{m} \sum_{\ell=1}^m I(a, b; \mathbf{X}^{(\ell)}, \mathbf{Y}^{(\ell)}, \mathbf{y}, f), \quad (1.7)$$

where m is the number of samples, $\mathbf{X}^{(\ell)} = \{X_1^{(\ell)}, \dots, X_n^{(\ell)}\}$, $\mathbf{Y}^{(\ell)} = \{Y_1^{(\ell)}, \dots, Y_n^{(\ell)}\}$, and $\{(X_i^{(\ell)}, Y_i^{(\ell)})\}$ is an iid sample with density $g(\cdot, \cdot)$.

In Section 2 we derive several “reference rules” for the kernel conditional density estimator making various assumption about the conditional density $f(y|x)$ and the marginal density $h(x)$.

In Section 3 we discuss an approximate parametric bootstrap method for estimating bandwidths, similar to that used by Hall et al. (1999) for bandwidth selection in estimating conditional distribution functions.

A third approach is considered in Section 4, where the estimation problem is written as a regression problem so that a bandwidth selection method from kernel regression can be modified for use here.

These various approaches to bandwidth selection are combined in Section 5 to provide a practical strategy for bandwidth selection. The methods are illustrated in Section 6 using two simulated examples and one real data set. Finally, we discuss extending the bandwidth selection methods to other estimators in Section 7.

2. Reference rules

Bandwidth rules based on a reference distribution have proven useful in univariate kernel density estimation (e.g., Silverman, 1986). The most common approach is to assume the underlying density is normal and find the bandwidth which would minimize the IMSE given that assumption. This is surprisingly robust and gives reasonable results even for densities which are quite non-normal. We shall apply the idea here to obtain a quick and simple method for bandwidth selection for kernel conditional density estimators.

Hyndman et al. (1996) showed the asymptotic mean square error for the estimator $\hat{f}(y|x)$ is

$$\begin{aligned} \text{AMSE} &= \lim_{n \rightarrow \infty} E\{\hat{f}(y|x) - f(y|x)\}^2 \\ &= \frac{a^4 \sigma_K^4}{4} \left\{ 2 \frac{h'(x)}{h(x)} \frac{\partial f(y|x)}{\partial x} + \frac{\partial^2 f(y|x)}{\partial x^2} + \frac{b^2}{a^2} \frac{\partial^2 f(y|x)}{\partial y^2} \right\}^2 \\ &\quad + \frac{R(K)f(y|x)}{nab h(x)} [R(K) - bf(y|x)] + O\left(\frac{1}{n}\right) + O\left(\frac{b}{an}\right) + O\left(\frac{a}{bn}\right) \\ &\quad + O(a^6) + O(b^6) + O(a^2 b^4) + O(a^4 b^2), \end{aligned} \quad (2.1)$$

where $R(K) = \int K^2(w) dw$. Then they show that substituting (2.1) into (1.4) gives

$$\text{IMSE} \approx \frac{c_1}{nab} - \frac{c_2}{na} + c_3 a^4 + c_4 b^4 + c_5 a^2 b^2, \quad (2.2)$$

where the constants c_1, c_2, c_3, c_4 and c_5 are given by

$$\begin{aligned} c_1 &= \int R^2(K) dx \\ c_2 &= \int \int R(K) f^2(y|x) dy dx \\ c_3 &= \int \int \frac{\sigma_K^4 h(x)}{4} \left\{ 2 \frac{h'(x)}{h(x)} \frac{\partial f(y|x)}{\partial x} + \frac{\partial^2 f(y|x)}{\partial x^2} \right\}^2 dy dx \\ c_4 &= \int \int \frac{\sigma_K^4 h(x)}{4} \left\{ \frac{\partial^2 f(y|x)}{\partial y^2} \right\}^2 dy dx \\ c_5 &= \int \int \frac{\sigma_K^4 h(x)}{2} \left\{ 2 \frac{h'(x)}{h(x)} \frac{\partial f(y|x)}{\partial x} + \frac{\partial^2 f(y|x)}{\partial x^2} \right\} \left\{ \frac{\partial^2 f(y|x)}{\partial y^2} \right\} dy dx \end{aligned}$$

and where the integrals are over the sample spaces of Y and X .

The optimal bandwidths can be derived by differentiating (2.2) with respect to a and b and setting the derivatives to zero. We require the sample space of X to be finite to ensure c_1 remains finite.

Hyndman et al. (1996) showed that the optimal bandwidths are approximately

$$a^* = c_1^{1/6} \left\{ 4 \left(\frac{c_3}{c_4} \right)^{1/4} + 2c_5 \left(\frac{c_3}{c_4} \right)^{3/4} \right\}^{-1/6} n^{-1/6} \quad (2.3)$$

and

$$b^* = a^* \left(\frac{c_3}{c_4} \right)^{1/4} = c_1^{1/6} \left\{ 4 \left(\frac{c_4^5}{c_3} \right)^{1/4} + 2c_5 \left(\frac{c_4}{c_3} \right)^{3/4} \right\}^{-1/6} n^{-1/6}. \quad (2.4)$$

We shall assume that the conditional distribution, $f(y|x)$, is normal with linear mean $r(x) = c + dx$ and linear standard deviation $\sigma(x) = p + qx$. Hence, $[Y|X=x] \stackrel{d}{=} N(c + dx, (p + qx)^2)$ and the conditional density is

$$f(y|x) = \frac{1}{(p + qx)\sqrt{2\pi}} \exp \left\{ \frac{-1}{2(p + qx)^2} (y - c - dx)^2 \right\}.$$

We shall substitute this expression for $f(y|x)$ into (2.3) and (2.4) to obtain reference rules for bandwidth selection. We also need to specify the marginal density $h(x)$. We consider two possibilities: a uniform marginal density over the space $[\ell, u]$ and a truncated normal marginal density with mean μ_h and standard deviation σ_h . The parameters of the assumed conditional and marginal distributions will be estimated from the data. When $q \neq 0$, we use iteratively reweighted least squares to estimate c and d , with p and q estimated by minimizing

$$\sum_{i=0}^n \{(y_i - \hat{c} - \hat{d}x_i)^2 - (p + qx_i)^2\}^2.$$

To differentiate between the various assumptions about the marginal distributions we will use the following notation for the reference rules:

Method	Marginal density $h(x)$	a value	b value
Reference rule U	Uniform	a_U	b_U
Reference rule N	Normal	a_N	b_N

For example a_U denotes the optimal value of a assuming $h(x)$ is a uniform density. The conditional distribution is always assumed to be normal.

The derivation of each reference rule requires extensive algebraic manipulation. The following rules were obtained with some help from the computer algebra package Mathematica (1999).

2.1. Uniform marginal distribution

To evaluate the reference rule for a and b with $h(x)$ uniform over $[\ell, u]$, we substitute the conditional and marginal densities into the constants c_1, \dots, c_5 . We then integrate the constants initially with respect to y over the sample space $(-\infty, \infty)$ and secondly with respect to x over the sample space $[\ell, u]$. The constant terms are

$$c_1 = R^2(K)(u - \ell), \quad c_2 = \frac{R(K)}{2q\sqrt{\pi}} \log \left(\frac{p + qu}{p + q\ell} \right), \quad c_3 = \frac{3}{512} \frac{\sigma_K^4 zw}{q\sqrt{\pi}(u - \ell)},$$

$$c_4 = \frac{3}{128} \frac{\sigma_K^4 z}{q\sqrt{\pi}(u-\ell)}, \quad c_5 = \frac{3}{128} \frac{\sigma_K^4 z(2d^2 - 3q^2)}{q\sqrt{\pi}(u-\ell)},$$

where $z = ((p+qu)^4 - (p+q\ell)^4) / ((p+qu)^4(p+q\ell)^4)$, $w = 19q^4 + 4d^4 + 28q^2d^2$, $d \neq 0$ and $q \neq 0$.

These values are then substituted into (2.3) and (2.4) to obtain the following reference rule:

$$a_U = \left\{ \frac{2^{15/2} \sqrt{\pi} R^2(K)(u-l)^2 q}{3n\sigma_K^4 z w^{3/4} [\sqrt{w} + 2d^2 - 3q^2]} \right\}^{1/6} \quad \text{and} \quad b_U = \frac{w^{1/4}}{\sqrt{2}} a_U. \quad (2.5)$$

Clearly the sizes of the bandwidths are affected by the assumptions made on the conditional density, marginal density and the number of observations.

Suppose we now assume that the conditional standard deviation is constant (let $q = 0$). We again substitute the conditional and marginal densities into the constant terms and obtain

$$c_1 = R^2(K)(u-l), \quad c_2 = \frac{R(K)(u-l)}{2p\sqrt{\pi}}, \quad c_3 = \frac{3}{32} \frac{\sigma_K^4 d^4}{p^5 \sqrt{\pi}},$$

$$c_4 = \frac{3}{32} \frac{\sigma_K^4}{p^5 \sqrt{\pi}}, \quad c_5 = \frac{3}{16} \frac{\sigma_K^4 d^2}{p^5 \sqrt{\pi}}.$$

This gives the following special case of the reference rule for $q = 0$:

$$a_U = \left\{ \frac{4\sqrt{\pi} R^2(K)(u-l)p^5}{3 n\sigma_K^4 d^5} \right\}^{1/6} \quad \text{and} \quad b_U = d a_U,$$

where $d \neq 0$. If $q = 0$ and $d = 0$ the conditional densities are equal for all x , and there is no need to condition on X .

2.2. Normal marginal distribution

We now assume that the marginal density $h(x)$ is normal with a constant mean μ_h and constant variance σ_h^2 . We further assume that conditional density $f(y|x)$ has a constant variance, that is $q = 0$. (We have not been able to solve the equations for the more general case of $q \neq 0$.) Following the same procedure as for the normal–uniform reference rule we initially integrate the constants over the sample space of y $(-\infty, \infty)$. Integrating the constants over the sample space of x $(-\infty, \infty)$ for a normal marginal density results in infinite bandwidths. Therefore, we choose limits of the integral over x to be $\mu_h \pm k\sigma_h$. Then we obtain

$$c_1 = 2k\sigma_h R^2(K), \quad c_2 = \frac{k\sigma_h R(K)}{p\sqrt{\pi}}, \quad c_3 = \frac{d^2 \sigma_K^4 v(k)}{32\pi\sqrt{2} p^5 \sigma_h^5},$$

$$c_4 = \frac{3\sigma_K^4 \lambda(k)}{32 p^5 \sqrt{\pi}}, \quad c_5 = \frac{3d^2 \sigma_K^4 \lambda(k)}{16 p^5 \sqrt{\pi}},$$

where $\lambda(k) = \int_{-k}^k \phi(t) dt$, $\phi(\cdot)$ is the standard normal density function, and

$$v(k) = \sqrt{2\pi} \sigma_h^3 (3d^2 \sigma_h^2 + 8p^2) \lambda(k) - 16k\sigma_h^2 p^2 e^{-k^2/2}.$$

Substituting the constants into (2.3) and (2.4) we obtain the following reference rule:

$$a_N = \left\{ \frac{16kR^2(K)p^5(288\pi^9\sigma_h^{58}\lambda^2(k))^{1/8}}{n\sigma_K^4 d^{5/2} v^{3/4}(k)[v^{1/2}(k) + d(18\pi\sigma_h^{10}\lambda^2(k))^{1/4}]} \right\}^{1/6},$$

$$b_N = \left\{ \frac{d^2 v(k)}{3\sqrt{2\pi}\sigma_h^5 \lambda(k)} \right\}^{1/4} a_N. \quad (2.6)$$

The value of k controls the size of the sample space in the x direction. Therefore as we increase k we also increase a_N and b_N . Common choices for k would be 2 or 3, and this would represent approximately 95% and 99.7% of the sample space, respectively.

3. A bootstrap bandwidth selection approach

Following the approach of Hall et al. (1999) for estimation of conditional distribution functions, we propose an approximate parametric bootstrap method. We fit a parametric model $Y_i = \beta_0 + \beta_1 X_i + \dots + \beta_k X_i^k + \sigma \varepsilon_i$ where ε_i are standard normal iid random variables, β_0, \dots, β_k and σ are estimated from the data and k is determined by Akaike's (1973) Information Criterion. We form a parametric estimator $\tilde{f}(y|x)$ based on the model. Then we simulate a bootstrap data set $\mathbf{Y}^{(\ell)} = \{Y_1^{(\ell)}, \dots, Y_n^{(\ell)}\}$ based on the observations $\mathbf{X} = \{X_1, \dots, X_n\}$. We choose a and b to minimize

$$\tilde{M}(a, b; m, \mathbf{y}', \tilde{f}) = \frac{1}{m} \sum_{\ell=1}^m I(a, b; \mathbf{X}, \mathbf{Y}^{(\ell)}, \mathbf{y}', \tilde{f}),$$

the bootstrap estimator of the IMSE (assuming the above parametric model).

Although this assumes the conditional density is unimodal, we have found it gives reasonable results even for multimodal densities, although (see Section 6) the procedure tends to result in overestimates of bandwidth b in the case of multimodal conditional densities.

This scheme is easily modified to other parametric models. For example, to allow for heteroscedasticity, replace σ by $(\sigma + vX_i)$ in the model. It would also be possible to fit a mixture of normals with polynomial mean to allow for multimodality.

4. A regression-based bandwidth selector

Fan et al. (1996) noted that the conditional density estimator $\hat{f}(y|x)$ obtained from (1.3) for given values of x and y is the value of β which minimizes the weighted least-squares function

$$\sum_{i=1}^n w_i(x) \{v_i(y) - \beta\}^2, \quad (4.1)$$

where $v_i(y) = b^{-1}K(|Y_i - y|/b)$ is a kernel function. For a given bandwidth b and a given value y , finding $\hat{f}(y|x)$ is a standard nonparametric problem of regressing

$v_i(y)$ on X_i . Fan, Yao and Tong use this idea to define local polynomial estimators of conditional densities. We shall exploit the idea by modifying a bandwidth selection method used in regression to derive an alternative method for selecting the bandwidth a given the bandwidth b . Härdle (1991) describes selecting the bandwidth for regression by minimizing the penalized average square prediction error.

For conditional density estimation, define the penalized average squared prediction error as

$$\begin{aligned} Q_b(a) &= \frac{\Delta}{n} \sum_{k=1}^N \sum_{i=1}^n \{v_i(y'_k) - \hat{f}(y'_k | X_i)\}^2 p(w_i(X_i)) \\ &= \frac{\Delta}{n} \sum_{k=1}^N \sum_{i=1}^n \left\{ v_i(y'_k) - \sum_{j=1}^n w_j(X_i) v_j(y'_k) \right\}^2 p(w_i(X_i)), \end{aligned}$$

where $\{y'_1, \dots, y'_N\}$ are equally spaced over the sample space of Y with $y_{i+1} - y_i = \Delta$, and where $p(u)$ is a penalty function with first order Taylor expansion $p(u) = 1 + 2u + O(u^2)$. In the numerical examples in Section 6, we use Akaike's (1974) finite prediction error $p(u) = (1 + u)/(1 - u)$.

Theorem 1. For fixed b , minimizing $Q_b(a)$ with respect to a is equivalent to minimizing the IMSE defined by (1.4).

We provide a proof of this result in the appendix.

For computational purposes, it is convenient to write $Q_b(a)$ as

$$Q_b(a) = \frac{\Delta}{n} \mathbf{p}^T (V - W^T V) \odot (V - W^T V) \mathbf{1},$$

where V is an $n \times N$ matrix with (i, j) th element $v_i(y'_j)$, W is an $n \times n$ matrix with (i, j) th element $w_j(x_i)$, \odot denotes the element-wise or Hadamard product, $\mathbf{1}$ denotes a vector of ones, and \mathbf{p} denotes the vector with i th element $p(w_i(x_i))$.

5. A practical bandwidth selection strategy

The preceding sections describe several bandwidth selection methods. The reference rules are fast and easily implemented but make strong assumptions about the data. The bootstrap method is less affected by the assumed distributions, but is slow to implement. The regression-based rule usually works well in finding a value for a , but it assumes b is given.

In this section, we describe an algorithm which effectively combines these methods to provide a practical bandwidth selection strategy:

1. Find an initial value for the smoothing parameter b using one of the reference rules. For most applications, we have found the rule with normal marginal density works well.
2. Given this value of b , use the regression-based method to find a value for a .

Table 1

Bandwidth estimates and IMSE values for Example 1. These are all means of 50 simulated samples each consisting of 100 observations

Method	a	b	IMSE ($\times 10^{-6}$)
Reference rule U ($q \neq 0$)	0.94	4.7	4.0
Reference rule U ($q = 0$)	0.98	4.9	3.8
Reference rule N ($k = 2$)	0.74	6.8	3.5
Reference rule N ($k = 3$)	0.80	7.4	3.6
Regress	0.93	7.5	4.0
Bootstrap	0.87	6.5	3.4
Combination	0.91	6.7	4.0
Optimal values	0.80	7.5	3.5

3. Use the bootstrap method to revise the estimate of b by minimizing $\tilde{M}(a, b; m, y', \tilde{f})$ with respect to b while holding a fixed at the value obtained in Step 2.

Steps 2 and 3 may be repeated one or more times. We have found this algorithm provides a relatively fast and useful approach to finding good bandwidths.

To illustrate the selection methods and the strategy described above, we shall use simulation on two examples and apply the methods to some real data.

6. Applications and comparisons

We compare the various bandwidth selection methods through two simulated models and by application to the Old Faithful Geyser data. In all cases, we have used the Gaussian kernel, $K(u) = \phi(u) = \exp(-u^2/2)/\sqrt{2\pi}$.

Example 1. Consider the simple model $Y_i = 10 + 5X_i + \varepsilon_i$ where $\{X_i\}$ and $\{\varepsilon_i\}$ are two independent sequences of normally distributed independent random variables with $X_1 \stackrel{d}{=} N(10, 9)$ and $\varepsilon_1 \stackrel{d}{=} N(0, 100)$. In this case, the optimal bandwidths are given by (2.6) as $a_N = 0.80$ and $b_N = 7.5$ (where $k = 3$).

We shall compare these with the estimated bandwidths obtained from the various methods. We shall also estimate the IMSE for each bandwidth selection method using $\hat{M}(\hat{a}, \hat{b}; m, y', f)$ from (1.7) with $m = 50$, the values of $\{y'_1, \dots, y'_N\}$ chosen to cover the interval $[-10, 130]$, $N = 25$, and $f(y|x) = \frac{1}{10}\phi((y - 10 - 5x)/10)$. For the bootstrap method, $m = 25$ is used in calculating $\tilde{M}(a, b; m, y', \tilde{f})$ for each a and b .

The bandwidths and estimated IMSE obtained are given in Table 1. These are the means of 50 simulated samples each consisting of $n = 100$ observations. Boxplots of the bandwidths and IMSE values are given in Figs. 4–6.

We note that the bootstrap selection method has resulted in a smaller IMSE than the optimal bandwidth choice; this is simply due to random variation in the samples selected. It is not surprising that the bootstrap and the N rule perform best as they




Fig. 4. Values of a for each method from 50 samples. The dotted line shows the optimal value of $a = 0.80$.




Fig. 5. Values of b for each method from 50 samples. The dotted line shows the optimal value of $b = 7.5$.

both assume the true underlying density in this case. Note that the b values for the regression method are obtained from the N rule. Both the bootstrap and combination methods tend to give lower values of b than the optimum. The main effect of using the combination method instead of the bootstrap method seems to be it increases the variability of the a value. However, it is much faster.




Fig. 6. Estimated ISE values for each method from 50 samples. The last boxplot shows ISE values for 50 samples using the optimal values of $a = 0.75$ and $b = 7.0$.




Fig. 7. Scatterplot of 100 observations from the model used in Example 2. Note the bimodality in the conditional densities for large X .

Example 2. In this example we use the model $Y_i = 2 \sin(\pi X_i) + \varepsilon_i$ where $\{X_i\}$ and $\{\varepsilon_i\}$ are two independent sequences of random variables with X_i uniformly distributed on $(0, 2)$ and $\varepsilon_i | X_i = W_i N_i + (1 - W_i) M_i$ where W_i is a binary variable with $\Pr(W_i = 1) = \Pr(W_i = 0) = 0.5$, $N_i \stackrel{d}{=} N(X_i, 0.09)$ and $M_i \stackrel{d}{=} N(0, 0.09)$. Fig. 7 shows a scatterplot of 100 observations from this model.

Table 2

Bandwidth estimates and IMSE values for Example 2. These are all means of 50 simulated samples each consisting of 100 observations

Method	a	b	IMSE ($\times 10^{-4}$)
Reference rule U ($q \neq 0$)	0.29	0.42	3.9
Reference rule U ($q = 0$)	0.32	0.45	4.1
Reference rule N ($k = 2$)	0.31	0.54	4.1
Reference rule N ($k = 3$)	0.31	0.58	4.2
Regress	0.059	0.57	2.6
Bootstrap	0.067	0.47	2.1
Combination	0.063	0.56	2.5
Optimal values	0.053	0.30	1.8

For this model, the optimal bandwidths can be found by minimizing the estimated IMSE $\hat{M}(a, b; m, y', f)$ where

$$f(y|x) = \frac{1}{0.6} \phi\left(\frac{y - 2 \sin(\pi x)}{0.3}\right) + \frac{1}{0.6} \phi\left(\frac{y - 2 \sin(\pi x) - x}{0.3}\right).$$

Using $m = 25$, $N = 25$, and the values of $\{y'_1, \dots, y'_N\}$ chosen to cover the interval $[-2.5, 2.5]$, we obtained optimal bandwidths of $a = 0.053$ and $b = 0.30$.

We shall also estimate the IMSE for each bandwidth selection method using (1.7). Again, we use $m = 50$ and the values of $\{y'_1, \dots, y'_N\}$ are chosen to cover the interval $[-2.5, 2.5]$ with $N = 25$. For the bootstrap method, $m = 25$ is used.

The bandwidths and estimated IMSE obtained are given in Table 2. As for Example 1, these are the means of 50 simulated samples each consisting of $n = 100$ observations. Boxplots of the bandwidths and IMSE values are given in Figs. 8–10. All the reference rule methods give values of a and b well above the optimum. Both the bimodality and nonlinear mean of the conditional distributions lead to smaller optimal bandwidths than under the assumptions behind the reference rules. However, the regression method is still selecting values of a close to the optimum despite assuming a value of b which is much too high. The combination and bootstrap methods both lead to good values for a in this case. However, b values from both methods are too large, probably because of the assumption of normality in the bootstrap procedure. As in Example 1, the combination method produces bandwidths with greater variability than the bootstrap method.

For all bandwidth selectors, there is an assumption of normality for the conditional density and this has led to the overestimation of b in each case. For a bimodal density, a smaller bandwidth is required than for a density as smooth as the normal density (see Silverman, 1986). The estimated densities computed using the values of a and b obtained from the bootstrap method (for example) still show the bimodality, but the peaks and troughs are not so sharply defined as they are when a smaller b is used.

This difficulty may be overcome using the pragmatic solution of Silverman of using 0.85 times the “optimal” b obtained from one of the proposed methods. This




Fig. 8. Values of a for each method from 50 samples. The dotted line shows the optimal value of $a = 0.75$.




Fig. 9. Values of b for each method from 50 samples. The dotted line shows the optimal value of $b = 7.0$.

approach still gives reasonably good bandwidths for unimodal densities, and performs better for bimodal densities. For the bootstrap selector, a less ad hoc solution is available by modifying the assumption of conditional normality to a conditional density comprising a mixture of normals.




Fig. 10. Estimated ISE values for each method from 50 samples. The last boxplot shows ISE values for 50 samples using the optimal values of $a = 0.05$ and $b = 0.30$.

Table 3
Bandwidth estimates for the old faithful Geyser data

Method	a	b
Reference rule U (q)	5.1	0.27
Reference rule U ($q = 0$)	6.1	0.33
Reference rule N ($k = 2$)	3.9	0.81
Reference rule N ($k = 3$)	4.1	0.87
Regress	2.2	0.87
Bootstrap	3.6	0.40
Combination	2.4	0.48

6.1. Old faithful Geyser data

Table 3 shows the results of applying the various bandwidth selectors to the data plotted in Fig. 1.

The conditional density estimator obtained using the bandwidth from the combination selector was shown in Fig. 2. We note that the visual impression of the conditional density estimator is relatively insensitive to small changes in the bandwidths. Any of the bandwidth selectors above results in a similar plot.

6.2. Computational efficiency

Consider the simple model $Y_i = 1 + X_i + \varepsilon_i$ where $\{X_i\}$ and $\{\varepsilon_i\}$ are iid standard normal random variables. We compare the computational time for each bandwidth selection method using 50 simulated samples with sample sizes of $n = 100, 500$ and

Table 4

Mean computation time for different methods. Times are measured in minutes:seconds

Rule	$n = 100$	$n = 500$	$n = 1000$
Normal reference rules	0:00	0:00	0:00
Regress	0:04	0:53	1:28
Bootstrap	5:16	12:45	22:05
Combination	1:41	5:19	7:45

1000. Using an IBM Thinkpad PC with a Pentium(R) II 300 MHz processor and 96M of RAM, the results were given in Table 4.

Clearly, the normal reference rules are much less time-consuming than other methods (one sample with $n = 10,000$ observations took 2 s). The bootstrap method was the most time consuming, with 1000 observations taking an average time of 22 min and 5 s.

7. Extensions to other estimators

Hyndman et al. (1996) considered a modified kernel estimator such that the new conditional density estimator had a mean function that could be specified by a smoother with better bias properties than that inherited by the standard kernel estimator, namely the Nadaraya–Watson smoother. They showed that the modified estimator had smaller IMSE under certain conditions.

Following the same approach as in Section 2, we find that the IMSE and the optimal bandwidths a^* and b^* of the modified kernel conditional density estimator take the same form as for the standard kernel conditional density estimator, except that the constants c_3 and c_5 are different:

$$c_3 = \int \int \frac{\sigma_K^4 h(x)}{4} \left\{ 2 \frac{h'(x)}{h(x)} \frac{\partial f(y - r(x) | x)}{\partial x} + \frac{\partial^2 f(y - r(x) | x)}{\partial x^2} \right\}^2 dy dx,$$

$$c_5 = \int \int \frac{\sigma_K^4 h(x)}{2} \left\{ 2 \frac{h'(x)}{h(x)} \frac{\partial f(y - r(x) | x)}{\partial x} + \frac{\partial^2 f(y - r(x) | x)}{\partial x^2} \right\} \left\{ \frac{\partial^2 f(y | x)}{\partial y^2} \right\} dy dx,$$

where $r(x) = E(Y | X = x)$ is the conditional mean.

Thus, reference rules can be derived for this estimator in the same way as for the standard estimator. We give just one example, the U rule with $q \neq 0$:

$$a_U = \left\{ \frac{2^{15/2} \sqrt{\pi} R^2(K) (u - l)^2}{3(19)^{3/4} n \sigma_K^4 z q^4 (\sqrt{19} - 3)} \right\}^{1/6} \quad \text{and} \quad b_U = (19/4)^{1/4} q a_U,$$

where z and w are defined as for (2.5). Note that this is the same as setting $d = 0$ in (2.5).

To extend these reference rules to Fan, Yao and Tong's local polynomial estimator, one would first need to derive the IMSE of that estimator using Theorem 1 of their

paper, then find expressions for the optimal values of a and b , analogous to (2.3) and (2.4). Extension of the reference rules to the case where there is a multivariate explanatory variable is more difficult.

The bootstrap selector can be easily applied to any estimator. The regression-based selector can be adapted to other estimators (including the multivariate case) by replacing $w_j(X_i)$ by the weight from the “equivalent kernel” obtained when the estimator is written as a linear smoother.

8. Conclusions

We have found conditional density estimation to be a useful data analytic tool for uncovering complex relationships between two or more variables (Hyndman et al., 1996). However, the use of this tool in practice has been hampered by the lack of a suitable bandwidth selection procedure.

In this paper, we have presented several bandwidth selection strategies for kernel conditional density estimation. The best performing strategy seems to be the bootstrap method outlined in Section 3. However, it is very slow and on current computing equipment, is only viable for small-to-medium size data sets.

When it is not possible to use the bootstrap method, we recommend using the practical bandwidth selection strategy (the combination estimator) outlined in Section 5.

Acknowledgements

Part of this work was carried out while Rob Hyndman was a visitor to the Department of Statistics, Colorado State University. Rob Hyndman was supported in part by an Australian Research Council grant.

Appendix. Proof of Theorem 1

Define

$$Q_b(a, y) = n^{-1} \sum_{i=1}^n \left\{ v_i(y) - \sum_{j=1}^n w_j(X_i) v_j(y) \right\}^2 p(w_i(X_i)) \quad (\text{A.1})$$

and note that

$$Q_b(a) = \Delta \sum_{k=1}^N Q_b(a, y'_k).$$

Substituting $\varepsilon_i = v_i(y) - f(y|X_i)$ into (A.1) and expanding the penalty term $p(w_i(X_i))$, we find

$$Q_b(a, y) \approx n^{-1} \sum_{i=1}^n \{ \varepsilon_i + f(y|X_i) - \hat{f}(y|X_i) \}^2 [1 + 2w_i(X_i)]. \quad (\text{A.2})$$

Expanding $Q_b(a, y)$ to find the leading terms and ignoring the lower-order terms we obtain

$$Q_b(a, y) = n^{-1} \sum_{i=1}^n \varepsilon_i^2 + \text{ASE}(a, y) + 2n^{-1} \sum_{i=1}^n \varepsilon_i [f(y | X_i) - \hat{f}(y | X_i)] \\ + 2n^{-1} \sum_{i=1}^n \varepsilon_i^2 w_i(X_i) + O(a^{-2}n^{-2}) + O(a^3n^{-1}), \quad (\text{A.3})$$

where

$$\text{ASE}(a, y) = n^{-1} \sum_{i=1}^n \{f(y | X_i) - \hat{f}(y | X_i)\}^2 \quad (\text{A.4})$$

denotes the average squared error.

We now show that the third and fourth terms on the right-hand side of (A.3) cancel each other out.

First, we compute the conditional expectation of the third summand of $Q_b(a, y)$:

$$\begin{aligned} E \left[2n^{-1} \sum_{i=1}^n \varepsilon_i [f(y | X_i) - \hat{f}(y | X_i)] | X_1 \dots X_n \right] \\ = 2n^{-1} \sum_{i=1}^n E \left[\varepsilon_i [f(y | X_i) - \sum_{j=1}^n w_j(X_i) v_j(y)] | X_1 \dots X_n \right] \\ = 2n^{-1} \sum_{i=1}^n E \left[\varepsilon_i [f(y | X_i) - \sum_{j=1}^n w_j(X_i) (\varepsilon_j + f(y | x_j))] | X_1 \dots X_n \right] \\ = 2n^{-1} \sum_{i=1}^n E \left[\varepsilon_i [f(y | X_i) - \sum_{j=1}^n w_j(X_i) f(y | x_j)] | X_1 \dots X_n \right] \\ - 2n^{-1} \sum_{i=1}^n E \left[\varepsilon_i \sum_{j=1}^n w_j(X_i) \varepsilon_j | X_1 \dots X_n \right] \\ = 2n^{-1} \sum_{i=1}^n E[\varepsilon_i | X_1 \dots X_n] \left[f(y | X_i) - \sum_{j=1}^n w_j(X_i) f(y | x_j) \right] \\ - 2n^{-1} \sum_{i=1}^n \sum_{j=1}^n w_j(X_i) E[\varepsilon_i \varepsilon_j | X_1 \dots X_n]. \end{aligned}$$

Now, $\{\varepsilon_i\}$ are independent random variables with $E(\varepsilon_i | X_1, \dots, X_n) = O(b^2)$ and variance $\sigma^2(X_i)$. The conditional expectation becomes

$$E \left[2n^{-1} \sum_{i=1}^n \varepsilon_i [f(y | X_i) - \hat{f}(y | X_i)] | X_1 \dots X_n \right] - 2n^{-1} \sum_{i=1}^n w_i(X_i) \sigma^2(X_i) + O(b^2).$$

The conditional expectation of the fourth summand in (A.3) is

$$E \left[2n^{-1} \sum_{i=1}^n \varepsilon_i^2 w_i(X_i) | X_1 \dots X_n \right] = 2n^{-1} \sum_{i=1}^n w_i(X_i) \sigma^2(X_i).$$

Thus, the conditional expectation of the third summand is approximately equal to the negative of the conditional expectation of the fourth summand of $Q_b(a, y)$, so that

$$Q_b(a, y) = n^{-1} \sum_{i=1}^n \varepsilon_i^2 + \text{ASE}(a, y) + O(b^2) + O(a^{-2}n^{-2}) + O(a^3n^{-1}).$$

Therefore,

$$Q_b(a) = \Delta \sum_{i=1}^n \varepsilon_i^2 + \frac{\Delta}{N} \sum_{k=1}^N \text{ASE}(a, y'_k) + O(b^2) + O(a^{-2}n^{-2}) + O(a^3n^{-1}).$$

Note that the first term in this expression is independent of a and that the second term is asymptotically equal to the IMSE defined by (1.4). Therefore, for fixed b , minimizing $Q_b(a)$ is asymptotically equivalent to minimizing the IMSE.

References

- Akaike, H., 1973. Information theory and an extension of the maximum likelihood principle. In: Petrov, B.N., Csaki, F. (Eds.), *Second International Symposium on Information Theory*. Adademiai Kiado, Budapest, pp. 267–281.
- Akaike, H., 1974. A new look at statistical model identification. *IEEE Trans. Automat. Control* AC 19, 716–723.
- Azzalini, A., Bowman, A.W., 1990. A look at some data on the Old Faithful geyser. *Appl. Statist.* 39, 357–365.
- Fan, J., Yao, Q., Tong, H., 1996. Estimation of conditional densities and sensitivity measures in nonlinear dynamical systems. *Biometrika* 83, 189–206.
- Hall, P., Wolff, R.C., Yao, Q., 1999. Methods for estimating a conditional distribution function. *J. Amer. Statist. Assoc.* 94, 154–163.
- Härdle, W., 1991. *Smoothing Techniques with Implementation in S*. Springer, New York.
- Hyndman, R.J., Bashtannyk, D.M., Grunwald, G.K., 1996. Estimating and visualizing conditional densities. *J. Comput. Graph. Statist.* 5, 315–336.
- Hyndman, R.J., Yao, Q., 1998. Nonparametric estimation and symmetry tests for conditional density functions. Working paper 17/98, Department of Econometrics and Business Statistics, Monash University.
- Mathematica 4, 1999. Wolfram Research Inc., Champaign, IL (www.mathematica.com).
- Rosenblatt, M., 1969. Conditional probability density and regression estimators. In: Krishnaiah, P.R. (Ed.), *Multivariate Analysis II*. Academic Press, New York, pp. 25–31.
- Silverman, B.W., 1986. *Density Estimation for Statistics and Data Analysis*. Chapman & Hall, London.
- S-Plus 2000 for Windows, 2000. MathSoft Inc., Cambridge MA (www.splus.mathsoft.com).
- Stone, C.J., 1994. The use of polynomial splines and their tensor products in multivariate function estimation. *Ann. Statist.* 22, 1994.
- Wand, M.P., Jones, M.C., 1995. *Kernel Smoothing*. Chapman & Hall, London.